

# GPT-2 Fine-tuning Pipeline for Next Word Prediction:

## A Comprehensive Code Analysis

Technical Documentation

July 9, 2025

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Overview . . . . .	3
1.2	Prerequisites . . . . .	3
<b>2</b>	<b>Theoretical Background</b>	<b>3</b>
2.1	Language Modeling . . . . .	3
2.2	Transformer Architecture Fundamentals . . . . .	3
2.2.1	Multi-Head Attention Mechanism . . . . .	3
2.2.2	Position Embeddings . . . . .	4
2.3	GPT-2 Architecture and Autoregressive Attention . . . . .	4
2.3.1	Key Architectural Differences . . . . .	4
2.3.2	Mathematical Formulation of GPT-2 Block . . . . .	4
2.3.3	Autoregressive vs Standard Multi-Head Attention . . . . .	5
2.3.4	Training Objective . . . . .	5
2.3.5	Attention Pattern Analysis . . . . .	6
2.4	Fine-tuning Process . . . . .	6
<b>3</b>	<b>Mathematical Implementation Details</b>	<b>6</b>
3.1	Complete GPT-2 Forward Pass . . . . .	6
3.2	Attention Mechanism Deep Dive . . . . .	7
3.2.1	Query, Key, Value Computation . . . . .	7
3.2.2	Causal Attention Computation . . . . .	7
3.2.3	Multi-Head Combination . . . . .	7
3.3	MLP Layer Mathematics . . . . .	7
3.4	Layer Normalization . . . . .	8
3.5	Training Loss and Optimization . . . . .	8
3.5.1	Cross-Entropy Loss . . . . .	8
3.5.2	Gradient Computation . . . . .	8
3.6	Comparison: Standard vs Causal Attention . . . . .	8
3.7	Implementation in PyTorch . . . . .	8
3.8	Import Statements . . . . .	10

<b>4</b>	<b>Dataset Implementation</b>	<b>10</b>
4.1	WikiTextDataset Class . . . . .	10
4.1.1	Key Features . . . . .	11
4.1.2	Data Flow . . . . .	11
4.2	Data Loading Functions . . . . .	11
<b>5</b>	<b>Model Setup and Configuration</b>	<b>12</b>
5.1	Model and Tokenizer Initialization . . . . .	12
5.1.1	Key Configurations . . . . .	12
<b>6</b>	<b>Training Pipeline</b>	<b>12</b>
6.1	Fine-tuning Function . . . . .	12
6.2	Training Arguments Analysis . . . . .	13
6.3	Data Collator . . . . .	13
<b>7</b>	<b>Evaluation Framework</b>	<b>13</b>
7.1	NWPEvaluator Class . . . . .	13
7.2	Perplexity Calculation . . . . .	14
7.3	Top-k Accuracy . . . . .	14
<b>8</b>	<b>Data Exploration Tools</b>	<b>15</b>
8.1	Training Data Inspection . . . . .	15
8.1.1	inspect_training_examples() . . . . .	15
8.1.2	show_next_word_prediction_examples() . . . . .	15
8.2	Statistical Analysis . . . . .	16
<b>9</b>	<b>Main Execution Pipeline</b>	<b>16</b>
9.1	Main Function . . . . .	16
9.2	Results Summary . . . . .	16
<b>10</b>	<b>Advanced Topics</b>	<b>17</b>
10.1	Attention Visualization . . . . .	17
10.2	Temperature and Sampling Strategies . . . . .	17
10.3	Memory-Efficient Training . . . . .	18
10.4	Distributed Training Considerations . . . . .	18
10.5	Basic Usage . . . . .	19
10.6	Custom Configuration . . . . .	19
<b>11</b>	<b>Performance Considerations</b>	<b>19</b>
11.1	Memory Management . . . . .	19
11.2	Computational Requirements . . . . .	20
<b>12</b>	<b>Best Practices and Recommendations</b>	<b>20</b>
12.1	Training Tips . . . . .	20
12.2	Evaluation Guidelines . . . . .	20
<b>13</b>	<b>Conclusion</b>	<b>20</b>
13.1	Future Enhancements . . . . .	21
<b>14</b>	<b>References</b>	<b>21</b>

# 1 Introduction

This document provides a comprehensive explanation of a GPT-2 fine-tuning pipeline designed for next word prediction tasks. The code implements a complete machine learning workflow from data preprocessing to model evaluation using PyTorch and Hugging Face Transformers.

## 1.1 Overview

The pipeline consists of several key components:

- Custom dataset handling for WikiText-2
- GPT-2 model fine-tuning using Hugging Face Transformers
- Comprehensive evaluation metrics (perplexity, top-k accuracy)
- Data exploration and visualization tools

## 1.2 Prerequisites

- Python 3.7+
- PyTorch
- Hugging Face Transformers
- Standard ML libraries (numpy, pandas, matplotlib)

# 2 Theoretical Background

## 2.1 Language Modeling

Language modeling is the task of predicting the next word in a sequence given the previous words. Formally, given a sequence of words  $w_1, w_2, \dots, w_{t-1}$ , the model predicts the probability distribution over the vocabulary for the next word  $w_t$ :

$$P(w_t | w_1, w_2, \dots, w_{t-1})$$

## 2.2 Transformer Architecture Fundamentals

### 2.2.1 Multi-Head Attention Mechanism

The core of the Transformer architecture is the multi-head attention mechanism. Given input sequences, attention computes relationships between all positions simultaneously.

**Scaled Dot-Product Attention:**  $\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$

where:

- $Q \in \mathbb{R}^{n \times d_k}$  is the query matrix
- $K \in \mathbb{R}^{n \times d_k}$  is the key matrix

- $V \in \mathbb{R}^{n \times d_v}$  is the value matrix
- $d_k$  is the dimension of keys/queries
- $n$  is the sequence length

**Multi-Head Attention:**  $\text{MultiHead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$   
 where each head is computed as:  $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$   
 The projection matrices are:

$$W_i^Q \in \mathbb{R}^{d_{\text{model}} \times d_k} \quad (1)$$

$$W_i^K \in \mathbb{R}^{d_{\text{model}} \times d_k} \quad (2)$$

$$W_i^V \in \mathbb{R}^{d_{\text{model}} \times d_v} \quad (3)$$

$$W^O \in \mathbb{R}^{hd_v \times d_{\text{model}}} \quad (4)$$

### 2.2.2 Position Embeddings

Since attention has no inherent notion of order, position embeddings are added:  $\text{PE}(\text{pos}, 2i) = \sin\left(\frac{\text{pos}}{10000^{2i/d_{\text{model}}}}\right)$   $\text{PE}(\text{pos}, 2i+1) = \cos\left(\frac{\text{pos}}{10000^{2i/d_{\text{model}}}}\right)$

## 2.3 GPT-2 Architecture and Autoregressive Attention

### 2.3.1 Key Architectural Differences

GPT-2 modifies the standard Transformer architecture for autoregressive generation:

**1. Causal (Masked) Self-Attention:** Unlike bidirectional attention in BERT, GPT-2 uses causal attention where each position can only attend to previous positions:

$$\text{Attention}_{\text{causal}}(Q, K, V) = \text{softmax}\left(\frac{QK^T + M}{\sqrt{d_k}}\right)V$$

where  $M$  is a lower triangular mask matrix:  $M_{ij} = \begin{cases} 0 & \text{if } i \geq j \\ -\infty & \text{if } i < j \end{cases}$

### 2. Architectural Components:

- **Decoder-only structure:** Unlike encoder-decoder models
- **Masked multi-head self-attention:** Prevents information leakage
- **Position embeddings:** Learned position embeddings (not sinusoidal)
- **Layer normalization:** Applied before attention and feedforward
- **Residual connections:** Skip connections around each sub-layer

### 2.3.2 Mathematical Formulation of GPT-2 Block

A single GPT-2 transformer block performs:

Attention Output:  $a = \text{LayerNorm}(x) + \text{CausalMultiHead}(\text{LayerNorm}(x))$  Block Output:  $y = a + \text{MLP}(\text{LayerNorm}(a))$

where the MLP is:  $\text{MLP}(x) = \text{GELU}(xW_1 + b_1)W_2 + b_2$

### 2.3.3 Autoregressive vs Standard Multi-Head Attention

#### Standard Multi-Head Attention (Bidirectional):

- Each position attends to ALL positions in the sequence
- Attention matrix is full:  $A_{ij}$  can be non-zero for all  $i, j$
- Used in BERT-style models for understanding tasks
- Attention weights:  $\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^n \exp(e_{ik})}$

#### Autoregressive (Causal) Attention:

- Each position only attends to previous positions (including itself)
- Attention matrix is lower triangular:  $A_{ij} = 0$  if  $i < j$
- Used in GPT-style models for generation tasks
- Attention weights:  $\alpha_{ij} = \frac{\exp(e_{ij})}{\sum_{k=1}^i \exp(e_{ik})}$  for  $j \leq i$

#### Implementation Comparison:

```

1 # Standard Multi-Head Attention
2 def standard_attention(Q, K, V):
3     scores = torch.matmul(Q, K.transpose(-2, -1)) / sqrt(d_k)
4     attn_weights = torch.softmax(scores, dim=-1)
5     output = torch.matmul(attn_weights, V)
6     return output
7
8 # Causal (GPT-2) Attention
9 def causal_attention(Q, K, V):
10    scores = torch.matmul(Q, K.transpose(-2, -1)) / sqrt(d_k)
11
12    # Apply causal mask
13    seq_len = scores.size(-1)
14    mask = torch.triu(torch.ones(seq_len, seq_len), diagonal=1)
15    mask = mask.masked_fill(mask == 1, float('-inf'))
16
17    scores = scores + mask
18    attn_weights = torch.softmax(scores, dim=-1)
19    output = torch.matmul(attn_weights, V)
20    return output

```

### 2.3.4 Training Objective

The model is trained to maximize the likelihood of the training data using teacher forcing:

$$\mathcal{L} = -\sum_{i=1}^N \log P(w_i | w_{<i}; \theta)$$

where  $\theta$  represents the model parameters and  $w_{<i}$  denotes all tokens before position  $i$ .

### 2.3.5 Attention Pattern Analysis

The attention patterns differ significantly:

$$\begin{aligned} \text{Bidirectional Attention Matrix: } A_{\text{bidirectional}} &= \begin{pmatrix} \alpha_{11} & \alpha_{12} & \alpha_{13} & \cdots & \alpha_{1n} \\ \alpha_{21} & \alpha_{22} & \alpha_{23} & \cdots & \alpha_{2n} \\ \alpha_{31} & \alpha_{32} & \alpha_{33} & \cdots & \alpha_{3n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \alpha_{n1} & \alpha_{n2} & \alpha_{n3} & \cdots & \alpha_{nn} \end{pmatrix} \\ \text{Causal Attention Matrix: } A_{\text{causal}} &= \begin{pmatrix} \alpha_{11} & 0 & 0 & \cdots & 0 \\ \alpha_{21} & \alpha_{22} & 0 & \cdots & 0 \\ \alpha_{31} & \alpha_{32} & \alpha_{33} & \cdots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \alpha_{n1} & \alpha_{n2} & \alpha_{n3} & \cdots & \alpha_{nn} \end{pmatrix} \end{aligned}$$

## 2.4 Fine-tuning Process

Fine-tuning involves taking a pre-trained model and adapting it to a specific task or domain by continuing training on task-specific data with a smaller learning rate.

## 3 Mathematical Implementation Details

### 3.1 Complete GPT-2 Forward Pass

The complete mathematical formulation of a GPT-2 forward pass:

**Input Processing:**

$$\text{Input tokens: } \mathbf{x} = [x_1, x_2, \dots, x_n] \quad (5)$$

$$\text{Token embeddings: } \mathbf{E}_{\text{tok}} = \text{Embedding}(\mathbf{x}) \quad (6)$$

$$\text{Position embeddings: } \mathbf{E}_{\text{pos}} = \text{PositionEmbedding}([1, 2, \dots, n]) \quad (7)$$

$$\text{Initial hidden state: } \mathbf{h}_0 = \mathbf{E}_{\text{tok}} + \mathbf{E}_{\text{pos}} \quad (8)$$

**Transformer Block Operations:** For each layer  $l = 1, 2, \dots, L$ :

$$\text{Pre-attention norm: } \mathbf{h}_l^{(1)} = \text{LayerNorm}(\mathbf{h}_{l-1}) \quad (9)$$

$$\text{Multi-head attention: } \mathbf{a}_l = \text{CausalMultiHead}(\mathbf{h}_l^{(1)}) \quad (10)$$

$$\text{Post-attention residual: } \mathbf{h}_l^{(2)} = \mathbf{h}_{l-1} + \mathbf{a}_l \quad (11)$$

$$\text{Pre-MLP norm: } \mathbf{h}_l^{(3)} = \text{LayerNorm}(\mathbf{h}_l^{(2)}) \quad (12)$$

$$\text{MLP computation: } \mathbf{m}_l = \text{MLP}(\mathbf{h}_l^{(3)}) \quad (13)$$

$$\text{Post-MLP residual: } \mathbf{h}_l = \mathbf{h}_l^{(2)} + \mathbf{m}_l \quad (14)$$

**Output Generation:**

$$\text{Final normalization: } \mathbf{h}_{\text{final}} = \text{LayerNorm}(\mathbf{h}_L) \quad (15)$$

$$\text{Logits: } \mathbf{logits} = \mathbf{h}_{\text{final}} \mathbf{W}_{\text{lm\_head}} \quad (16)$$

$$\text{Probabilities: } \mathbf{p} = \text{softmax}(\mathbf{logits}) \quad (17)$$

## 3.2 Attention Mechanism Deep Dive

### 3.2.1 Query, Key, Value Computation

For each attention head  $i$  in layer  $l$ :

$$\mathbf{Q}_i = \mathbf{h}_l^{(1)} \mathbf{W}_i^Q + \mathbf{b}_i^Q \quad (18)$$

$$\mathbf{K}_i = \mathbf{h}_l^{(1)} \mathbf{W}_i^K + \mathbf{b}_i^K \quad (19)$$

$$\mathbf{V}_i = \mathbf{h}_l^{(1)} \mathbf{W}_i^V + \mathbf{b}_i^V \quad (20)$$

where  $\mathbf{W}_i^Q, \mathbf{W}_i^K, \mathbf{W}_i^V \in \mathbb{R}^{d_{model} \times d_k}$  and  $d_k = d_{model}/h$ .

### 3.2.2 Causal Attention Computation

The causal attention for head  $i$  is computed as:

$$\mathbf{S}_i = \frac{\mathbf{Q}_i \mathbf{K}_i^T}{\sqrt{d_k}} \quad (21)$$

$$\mathbf{M}_{jk} = \begin{cases} 0 & \text{if } j \geq k \\ -\infty & \text{if } j < k \end{cases} \quad (22)$$

$$\mathbf{A}_i = \text{softmax}(\mathbf{S}_i + \mathbf{M}) \quad (23)$$

$$\mathbf{O}_i = \mathbf{A}_i \mathbf{V}_i \quad (24)$$

### 3.2.3 Multi-Head Combination

The outputs from all heads are concatenated and projected:

$$\mathbf{O}_{\text{concat}} = \text{Concat}(\mathbf{O}_1, \mathbf{O}_2, \dots, \mathbf{O}_h) \quad (25)$$

$$\mathbf{O}_{\text{final}} = \mathbf{O}_{\text{concat}} \mathbf{W}^O + \mathbf{b}^O \quad (26)$$

## 3.3 MLP Layer Mathematics

The MLP in each transformer block:

$$\mathbf{h}_{\text{intermediate}} = \text{GELU}(\mathbf{h}_l^{(3)} \mathbf{W}_1 + \mathbf{b}_1) \quad (27)$$

$$\mathbf{h}_{\text{output}} = \mathbf{h}_{\text{intermediate}} \mathbf{W}_2 + \mathbf{b}_2 \quad (28)$$

where  $\mathbf{W}_1 \in \mathbb{R}^{d_{model} \times 4d_{model}}$  and  $\mathbf{W}_2 \in \mathbb{R}^{4d_{model} \times d_{model}}$ .

**GELU Activation:**  $\text{GELU}(x) = x \cdot \Phi(x) = x \cdot \frac{1}{2} \left[ 1 + \text{erf} \left( \frac{x}{\sqrt{2}} \right) \right]$

### 3.4 Layer Normalization

GPT-2 uses layer normalization before each sub-layer:

$$\mu = \frac{1}{d_{model}} \sum_{i=1}^{d_{model}} x_i \quad (29)$$

$$\sigma^2 = \frac{1}{d_{model}} \sum_{i=1}^{d_{model}} (x_i - \mu)^2 \quad (30)$$

$$\text{LayerNorm}(x) = \gamma \frac{x - \mu}{\sqrt{\sigma^2 + \epsilon}} + \beta \quad (31)$$

where  $\gamma$  and  $\beta$  are learnable parameters.

### 3.5 Training Loss and Optimization

#### 3.5.1 Cross-Entropy Loss

For next token prediction, the loss at each position is:

$$\mathcal{L}_i = -\log P(w_i | w_{<i}) = -\log \frac{\exp(\text{logits}_{i,w_i})}{\sum_{v \in V} \exp(\text{logits}_{i,v})}$$

The total loss is:  $\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \mathcal{L}_i$

#### 3.5.2 Gradient Computation

The gradient of the loss with respect to logits:

$$\frac{\partial \mathcal{L}}{\partial \text{logits}_{i,v}} = P(v | w_{<i}) - \mathbf{1}_{v=w_i}$$

where  $\mathbf{1}_{v=w_i}$  is 1 if  $v$  is the true token and 0 otherwise.

### 3.6 Comparison: Standard vs Causal Attention

Aspect	Standard (BERT)	Causal (GPT-2)
Attention Matrix	Full matrix	Lower triangular
Information Flow	Bidirectional	Unidirectional
Masking	Random tokens	Future tokens
Training Objective	MLM + NSP	Next token prediction
Use Case	Understanding	Generation
Computational Cost	$O(n^2)$	$O(n^2)$
Memory Pattern	Full attention	Causal attention

Table 1: Comparison of Attention Mechanisms

### 3.7 Implementation in PyTorch

Here's how the core attention mechanisms are implemented:



```

1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 import math
5
6 class CausalMultiHeadAttention(nn.Module):
7     def __init__(self, d_model, n_heads):
8         super().__init__()
9         self.d_model = d_model
10        self.n_heads = n_heads
11        self.d_k = d_model // n_heads
12
13        self.W_q = nn.Linear(d_model, d_model)
14        self.W_k = nn.Linear(d_model, d_model)
15        self.W_v = nn.Linear(d_model, d_model)
16        self.W_o = nn.Linear(d_model, d_model)
17
18        # Create causal mask
19        self.register_buffer('mask', torch.triu(torch.ones(512, 512),
20        diagonal=1))
21
22    def forward(self, x):
23        batch_size, seq_len, d_model = x.size()
24
25        # Compute Q, K, V
26        Q = self.W_q(x).view(batch_size, seq_len, self.n_heads, self.
27        d_k).transpose(1, 2)
28        K = self.W_k(x).view(batch_size, seq_len, self.n_heads, self.
29        d_k).transpose(1, 2)
30        V = self.W_v(x).view(batch_size, seq_len, self.n_heads, self.
31        d_k).transpose(1, 2)
32
33        # Compute attention scores
34        scores = torch.matmul(Q, K.transpose(-2, -1)) / math.sqrt(self.
35        d_k)
36
37        # Apply causal mask
38        mask = self.mask[:seq_len, :seq_len]
39        scores = scores.masked_fill(mask == 1, float('-inf'))
40
41        # Apply softmax
42        attn_weights = F.softmax(scores, dim=-1)
43
44        # Apply attention to values
45        out = torch.matmul(attn_weights, V)
46
47        # Concatenate heads
48        out = out.transpose(1, 2).contiguous().view(batch_size, seq_len
49        , d_model)
50
51        # Final linear transformation
52        out = self.W_o(out)
53
54        return out
55
56 class GPT2Block(nn.Module):
57     def __init__(self, d_model, n_heads, d_ff):
58         super().__init__()

```

```

53     self.ln_1 = nn.LayerNorm(d_model)
54     self.attn = CausalMultiHeadAttention(d_model, n_heads)
55     self.ln_2 = nn.LayerNorm(d_model)
56     self.mlp = nn.Sequential(
57         nn.Linear(d_model, d_ff),
58         nn.GELU(),
59         nn.Linear(d_ff, d_model)
60     )
61
62     def forward(self, x):
63         # Pre-norm architecture
64         x = x + self.attn(self.ln_1(x))
65         x = x + self.mlp(self.ln_2(x))
66         return x

```

### 3.8 Import Statements

The code begins with comprehensive imports:

```

1  import torch
2  import torch.nn as nn
3  from torch.utils.data import Dataset, DataLoader
4  from transformers import (
5      GPT2LMHeadModel, GPT2Tokenizer,
6      TrainingArguments, Trainer,
7      DataCollatorForLanguageModeling
8  )
9  from datasets import load_dataset
10 # ... additional imports

```

These imports provide:

- **torch**: Core PyTorch functionality
- **transformers**: Hugging Face’s transformer models
- **datasets**: Dataset loading utilities
- **Standard libraries**: For data processing and visualization

## 4 Dataset Implementation

### 4.1 WikiTextDataset Class

The WikiTextDataset class handles data preprocessing for the WikiText-2 dataset:

```

1  class WikiTextDataset(Dataset):
2      def __init__(self, texts: List[str], tokenizer, max_length: int =
3          512):
4          self.tokenizer = tokenizer
5          self.max_length = max_length
6          self.examples = []
7
8          for text in tqdm(texts):
9              if len(text.strip()) > 0:
10                 tokens = tokenizer.encode(text, add_special_tokens=True
11             )

```

```

10
11         # Create overlapping sequences
12         for i in range(0, len(tokens) - 1, max_length // 2):
13             chunk = tokens[i:i + max_length]
14             if len(chunk) > 1:
15                 self.examples.append(chunk)

```

#### 4.1.1 Key Features

1. **Tokenization:** Converts text to token IDs using GPT-2 tokenizer
2. **Sliding Window:** Creates overlapping sequences with 50% overlap
3. **Length Filtering:** Ensures sequences have minimum length for training

#### 4.1.2 Data Flow

---

##### Algorithm 1 Dataset Creation Process

---

```

for each text in input texts do
    if text is not empty then
        tokenize text to get token IDs
        for i = 0 to length(tokens) - 1 step max_length/2 do
            chunk = tokens[i:i + max_length]
            if length(chunk) > 1 then
                add chunk to examples
            end if
        end for
    end if
end for

```

---

## 4.2 Data Loading Functions

The `load_and_prepare_data()` function handles dataset loading:

```

1 def load_and_prepare_data():
2     dataset = load_dataset('wikitext', 'wikitext-2-raw-v1')
3
4     train_texts = [text for text in dataset['train']['text']
5                     if len(text.strip()) > 0]
6     valid_texts = [text for text in dataset['validation']['text']
7                    if len(text.strip()) > 0]
8     test_texts = [text for text in dataset['test']['text']
9                   if len(text.strip()) > 0]
10
11     return train_texts, valid_texts, test_texts

```

This function:

- Loads WikiText-2 dataset from Hugging Face
- Filters empty texts
- Returns train/validation/test splits

## 5 Model Setup and Configuration

### 5.1 Model and Tokenizer Initialization

The `setup_model_and_tokenizer()` function configures the GPT-2 components:

```
1 def setup_model_and_tokenizer():
2     model_name = "gpt2"
3
4     tokenizer = GPT2Tokenizer.from_pretrained(model_name)
5     model = GPT2LMHeadModel.from_pretrained(model_name)
6
7     # Add padding token
8     tokenizer.pad_token = tokenizer.eos_token
9
10    return model, tokenizer
```

#### 5.1.1 Key Configurations

- **Model:** GPT2LMHeadModel with language modeling head
- **Tokenizer:** GPT2Tokenizer for text-to-token conversion
- **Padding:** Uses EOS token as padding token

## 6 Training Pipeline

### 6.1 Fine-tuning Function

The `fine_tune_model()` function implements the training loop:

```
1 def fine_tune_model(model, tokenizer, train_dataset, eval_dataset,
2                     output_dir="./gpt2-wikitext2"):
3
4     data_collator = DataCollatorForLanguageModeling(
5         tokenizer=tokenizer,
6         mlm=False, # GPT-2 is autoregressive
7         pad_to_multiple_of=8,
8         return_tensors="pt"
9     )
10
11     training_args = TrainingArguments(
12         output_dir=output_dir,
13         num_train_epochs=3,
14         per_device_train_batch_size=4,
15         per_device_eval_batch_size=4,
16         warmup_steps=500,
17         learning_rate=5e-5,
18         weight_decay=0.01,
19         fp16=True,
20         # ... additional arguments
21     )
22
23     trainer = Trainer(
24         model=model,
```

```

25     args=training_args ,
26     data_collator=data_collator ,
27     train_dataset=train_dataset ,
28     eval_dataset=eval_dataset ,
29     tokenizer=tokenizer ,
30 )
31
32 trainer.train()
33 return trainer

```

## 6.2 Training Arguments Analysis

Parameter	Description
num_train_epochs	Number of training epochs (3)
per_device_train_batch_size	Training batch size per device (4)
warmup_steps	Learning rate warmup steps (500)
learning_rate	Initial learning rate (5e-5)
weight_decay	L2 regularization coefficient (0.01)
fp16	Mixed precision training (True)
gradient_accumulation_steps	Gradient accumulation (2)

Table 2: Key Training Parameters

## 6.3 Data Collator

The `DataCollatorForLanguageModeling` handles:

- Padding sequences to equal length
- Creating attention masks
- Preparing labels for next token prediction

# 7 Evaluation Framework

## 7.1 NWPEvaluator Class

The `NWPEvaluator` class provides comprehensive evaluation metrics:

```

1 class NWPEvaluator:
2     def __init__(self, model, tokenizer, device):
3         self.model = model
4         self.tokenizer = tokenizer
5         self.device = device
6         self.model.eval()
7
8     def calculate_perplexity(self, dataloader) -> float:
9         # Implementation for perplexity calculation
10
11     def calculate_top_k_accuracy(self, dataloader,

```

```

12         k_values: List[int] = [1, 5, 10]) ->
    Dict[int, float]:
13     # Implementation for top-k accuracy
14
15     def sample_predictions(self, text: str,
16                           num_predictions: int = 5) -> List[str]:
17     # Implementation for sample predictions

```

## 7.2 Perplexity Calculation

Perplexity measures how well the model predicts the next word:

$$\text{Perplexity} = \exp \left( -\frac{1}{N} \sum_{i=1}^N \log P(w_i | w_{<i}) \right)$$

Implementation:

```

1 def calculate_perplexity(self, dataloader) -> float:
2     total_loss = 0
3     total_tokens = 0
4
5     with torch.no_grad():
6         for batch in dataloader:
7             input_ids = batch['input_ids'].to(self.device)
8             attention_mask = batch['attention_mask'].to(self.device)
9
10            # Shift labels for next token prediction
11            labels = input_ids.clone()
12            labels[:, :-1] = input_ids[:, 1:]
13            labels[:, -1] = -100 # Ignore last token
14
15            outputs = self.model(input_ids=input_ids,
16                                attention_mask=attention_mask,
17                                labels=labels)
18
19            loss = outputs.loss
20            valid_tokens = (labels != -100).sum().item()
21
22            total_loss += loss.item() * valid_tokens
23            total_tokens += valid_tokens
24
25        avg_loss = total_loss / total_tokens
26        perplexity = math.exp(avg_loss)
27    return perplexity

```

## 7.3 Top-k Accuracy

Top-k accuracy measures if the correct next word appears in the top-k predictions:

```

1 def calculate_top_k_accuracy(self, dataloader, k_values: List[int] =
    [1, 5, 10]):
2     correct_predictions = {k: 0 for k in k_values}
3     total_predictions = 0
4
5     with torch.no_grad():
6         for batch in dataloader:

```

```

7         # Process each sequence
8         for seq_idx in range(input_ids.size(0)):
9             sequence = input_ids[seq_idx]
10
11         # Predict each position
12         for pos in range(1, valid_length):
13             context = sequence[:pos].unsqueeze(0)
14             target = sequence[pos].item()
15
16             outputs = self.model(context)
17             logits = outputs.logits[0, -1, :]
18
19             # Get top-k predictions
20             top_k_tokens = torch.topk(logits, max(k_values)).
indices
21
22             # Check accuracy for different k values
23             for k in k_values:
24                 if target in top_k_tokens[:k]:
25                     correct_predictions[k] += 1
26
27             total_predictions += 1
28
29         accuracies = {k: correct_predictions[k] / total_predictions for k
in k_values}
30         return accuracies

```

## 8 Data Exploration Tools

### 8.1 Training Data Inspection

The code includes comprehensive data exploration functions:

#### 8.1.1 inspect\_training\_examples()

This function provides detailed analysis of training examples:

- Sequence length statistics
- Token-level breakdown
- Text representation with special tokens
- Attention mask visualization

#### 8.1.2 show\_next\_word\_prediction\_examples()

Demonstrates the input-output format for next word prediction:

```

1 def show_next_word_prediction_examples(dataset, tokenizer, num_examples
=3):
2     for i in range(min(num_examples, len(dataset))):
3         example = dataset[i]
4         input_ids = example['input_ids']
5

```

```

6         # Show first 10 positions as input -> target pairs
7         for pos in range(1, min(11, len(input_ids))):
8             context_ids = input_ids[:pos]
9             target_id = input_ids[pos]
10
11             context_text = tokenizer.decode(context_ids,
12 skip_special_tokens=True)
13             target_text = tokenizer.decode([target_id],
14 skip_special_tokens=True)
15
16             print(f"'{context_text}' -> '{target_text}'")

```

## 8.2 Statistical Analysis

The `analyze_dataset_statistics()` function provides:

- Sequence length distribution
- Token frequency analysis
- Vocabulary statistics
- Data quality metrics

## 9 Main Execution Pipeline

### 9.1 Main Function

The `main()` function orchestrates the entire pipeline:

---

#### Algorithm 2 Main Execution Flow

---

```

Set device (GPU/CPU)
Load and prepare data
Setup model and tokenizer
Create datasets
Fine-tune model
Evaluate model
Generate sample predictions
Create results summary

```

---

### 9.2 Results Summary

The pipeline generates comprehensive results including:

- Perplexity scores
- Top-k accuracy metrics
- Sample predictions
- Performance visualizations



## 10 Advanced Topics

### 10.1 Attention Visualization

Understanding attention patterns is crucial for interpreting model behavior:

```

1 def visualize_attention_patterns(model, tokenizer, text):
2     """Visualize attention patterns for a given text"""
3     inputs = tokenizer(text, return_tensors="pt")
4
5     with torch.no_grad():
6         outputs = model(**inputs, output_attentions=True)
7         attentions = outputs.attentions # List of attention tensors
8
9     # attentions[i] has shape: [batch_size, num_heads, seq_len, seq_len]
10    # For layer i, head j: attentions[i][0, j, :, :]
11
12    import matplotlib.pyplot as plt
13
14    # Visualize attention for first layer, first head
15    attn_matrix = attentions[0][0, 0, :, :].numpy()
16
17    plt.figure(figsize=(10, 8))
18    plt.imshow(attn_matrix, cmap='Blues')
19    plt.colorbar()
20    plt.title("Causal Attention Pattern (Layer 0, Head 0)")
21    plt.xlabel("Key Position")
22    plt.ylabel("Query Position")
23    plt.show()

```

### 10.2 Temperature and Sampling Strategies

The code can be extended with different sampling strategies:

```

1 def sample_with_temperature(logits, temperature=1.0):
2     """Sample from logits with temperature scaling"""
3     if temperature == 0:
4         return torch.argmax(logits, dim=-1)
5
6     # Apply temperature scaling
7     scaled_logits = logits / temperature
8     probs = F.softmax(scaled_logits, dim=-1)
9
10    # Sample from the distribution
11    return torch.multinomial(probs, 1)
12
13 def top_k_top_p_sampling(logits, top_k=50, top_p=0.9, temperature=1.0):
14     """Combined top-k and top-p (nucleus) sampling"""
15     # Apply temperature
16     logits = logits / temperature
17
18     # Top-k filtering
19     if top_k > 0:
20         top_k_logits, top_k_indices = torch.topk(logits, top_k)
21         logits_filtered = torch.full_like(logits, float('-inf'))
22         logits_filtered.scatter_(1, top_k_indices, top_k_logits)

```

```

23     logits = logits_filtered
24
25     # Top-p filtering
26     if top_p < 1.0:
27         sorted_logits, sorted_indices = torch.sort(logits, descending=
True)
28         cumulative_probs = torch.cumsum(F.softmax(sorted_logits, dim
=-1), dim=-1)
29
30         # Remove tokens with cumulative probability above threshold
31         sorted_indices_to_remove = cumulative_probs > top_p
32         sorted_indices_to_remove[..., 1:] = sorted_indices_to_remove
[... , :-1].clone()
33         sorted_indices_to_remove[..., 0] = 0
34
35         indices_to_remove = sorted_indices_to_remove.scatter(1,
sorted_indices, sorted_indices_to_remove)
36         logits = logits.masked_fill(indices_to_remove, float('-inf'))
37
38     # Sample from filtered distribution
39     probs = F.softmax(logits, dim=-1)
40     return torch.multinomial(probs, 1)

```

## 10.3 Memory-Efficient Training

For large-scale training, memory optimization is crucial:

```

1 from torch.utils.checkpoint import checkpoint
2
3 class MemoryEfficientGPT2Block(nn.Module):
4     def __init__(self, d_model, n_heads, d_ff):
5         super().__init__()
6         self.ln_1 = nn.LayerNorm(d_model)
7         self.attn = CausalMultiHeadAttention(d_model, n_heads)
8         self.ln_2 = nn.LayerNorm(d_model)
9         self.mlp = nn.Sequential(
10             nn.Linear(d_model, d_ff),
11             nn.GELU(),
12             nn.Linear(d_ff, d_model)
13         )
14
15     def forward(self, x):
16         # Use gradient checkpointing to save memory
17         def attention_forward(x):
18             return self.attn(self.ln_1(x))
19
20         def mlp_forward(x):
21             return self.mlp(self.ln_2(x))
22
23         # Apply gradient checkpointing
24         x = x + checkpoint(attention_forward, x)
25         x = x + checkpoint(mlp_forward, x)
26         return x

```

## 10.4 Distributed Training Considerations

For multi-GPU training, the code can be extended:

```

1 import torch.distributed as dist
2 from torch.nn.parallel import DistributedDataParallel as DDP
3
4 def setup_distributed_training():
5     """Setup for distributed training across multiple GPUs"""
6     if torch.cuda.is_available() and torch.cuda.device_count() > 1:
7         dist.init_process_group(backend='nccl')
8         local_rank = int(os.environ['LOCAL_RANK'])
9         torch.cuda.set_device(local_rank)
10
11         # Wrap model with DDP
12         model = DDP(model, device_ids=[local_rank])
13
14         # Use DistributedSampler for data loading
15         sampler = torch.utils.data.distributed.DistributedSampler(
16             dataset)
17         dataloader = DataLoader(dataset, sampler=sampler, batch_size=
18             batch_size)
19
20         return model, dataloader
21
22 return model, dataloader

```

## 10.5 Basic Usage

To run the complete pipeline:

```

1 # Run full pipeline
2 model, tokenizer, results = main()
3
4 # Explore data first
5 sample_train, sample_valid, tokenizer = explore_training_data()

```

## 10.6 Custom Configuration

For custom training configurations:

```

1 # Custom dataset size
2 train_dataset = WikiTextDataset(train_texts[:5000], tokenizer)
3 eval_dataset = WikiTextDataset(valid_texts[:1000], tokenizer)
4
5 # Custom training arguments
6 training_args = TrainingArguments(
7     output_dir="./custom-gpt2",
8     num_train_epochs=5,
9     per_device_train_batch_size=8,
10    learning_rate=3e-5,
11    # ... other parameters
12 )

```

# 11 Performance Considerations

## 11.1 Memory Management

The code includes several memory optimization strategies:

- Mixed precision training (fp16)
- Gradient accumulation
- Efficient data loading
- Batch size optimization

## 11.2 Computational Requirements

Component	Requirement
GPU Memory	8GB+ recommended
Training Time	2-4 hours (depends on data size)
Disk Space	2GB+ for model and data
RAM	16GB+ recommended

Table 3: Computational Requirements

## 12 Best Practices and Recommendations

### 12.1 Training Tips

1. Start with small dataset subsets for testing
2. Monitor training loss and validation metrics
3. Use appropriate learning rates (typically  $1e-5$  to  $5e-5$ )
4. Implement early stopping for overfitting prevention

### 12.2 Evaluation Guidelines

1. Use multiple evaluation metrics
2. Test on held-out data
3. Compare with baseline models
4. Analyze failure cases

## 13 Conclusion

This GPT-2 fine-tuning pipeline provides a comprehensive framework for next word prediction tasks. The code implements best practices for:

- Data preprocessing and tokenization
- Model training and optimization

- Comprehensive evaluation
- Result analysis and visualization

The modular design allows for easy customization and extension for specific use cases while maintaining robust performance and reliability.

### 13.1 Future Enhancements

Potential improvements include:

- Distributed training support
- Advanced regularization techniques
- Real-time inference optimization
- Integration with larger language models

## 14 References

1. Radford, A., et al. (2019). Language models are unsupervised multitask learners.
2. Vaswani, A., et al. (2017). Attention is all you need. NIPS.
3. Wolf, T., et al. (2020). Transformers: State-of-the-art natural language processing.
4. Merity, S., et al. (2016). Pointer sentinel mixture models.